

Sample Document - PDF Semantic Search

This sample PDF demonstrates the ingestion and semantic search pipeline. The service extracts text from PDF files, splits the text into overlapping chunks, generates vector embeddings for each chunk with a sentence-transformer model, and stores them in the Qdrant vector database.

A search query is embedded the same way and compared against the stored vectors using cosine similarity, returning the most semantically relevant chunks, so you can find content by meaning rather than by exact keywords.